

Manuel Scionti

Barcelona, Spain | scionti.manuel@hotmail.it | +34 614 059 305
linkedin.com/in/manuelscionti | mnlsn.github.io | github.com/mnlsn

Professional Summary

AI Engineer focused on the part of AI that starts after the prototype: retrieval, agent orchestration, evaluation, and the infrastructure that makes LLM systems reliable in production. 3+ years across Fortune 500 enterprise (ADP) and AI startups, working across Databricks, AWS, and Azure to turn model capabilities into software teams can rely on.

Core Skills

LLM Systems and Retrieval: Large Language Models (LLMs), Generative AI, Retrieval-Augmented Generation (RAG), Graph RAG, agent workflows, multi-agent systems, LLM-as-judge, evaluation pipelines, semantic search, keyword search, vector search, OpenSearch, prompt engineering

Languages and APIs: Python, SQL, FastAPI, REST APIs, serverless APIs

Cloud and Platforms: AWS, AWS Lambda, API Gateway, SageMaker, Bedrock, Azure, Databricks, Hydra Compose, YAML configuration, CI/CD, Jenkins, Bitbucket, MLOps, observability

ML and Multimodal: NLP, computer vision, text-to-speech, speech-to-text, multimodal retrieval, clustering, customer segmentation

Professional Experience

AI Engineer

Mar 2025 – Present

ADP - Automatic Data Processing Inc. (via Capitale)

Barcelona, Spain

- Work on a team of 5 engineers designing and maintaining an AI platform for automated FAQ generation, built on an internal LLM orchestration framework and API layer
- Took 4 of 15+ supported business units from pilot to production; serving 1000+ daily users across US and EMEA regions
- Core contributor to an internal LLM orchestration framework used by multiple teams across the company; focused on reusability and code quality, cutting new pipeline integration development time by 50%
- Developed an evaluation suite (10+ evals) with scoring and a golden dataset to measure RAG quality; improved retrieval quality by 20% on the internal golden dataset through a reranking component and advanced prompting techniques
- Led an internal workshop (1 hour, 20 attendees) on leveraging AI coding agents, covering best practices and real-life examples; produced actionable workflows to automate and update documentation across the repo
- Refactored the Jenkins CI/CD pipeline to enable step modularity and parallelization (unit tests, integration, security scan, packaging, Artifactory push, docs); reduced build time from 25 to 16 minutes (36%)

AI Engineer

Nov 2023 – Mar 2025

Neodata Group

Catania, Italy

- Conducted feasibility studies for 20+ enterprise clients (all with €5M+ revenue), translating business needs into technical requirements with fast iteration cycles
- Led the design and prototyping of Neovid – a multimodal retrieval system over large image, video, and document collections (NLP, computer vision, face recognition, TTS, STT); processed 1000+ videos for Italy's largest broadcaster
- Main developer of NeoKnowledge: AI pipelines generating personalized educational content (text, audio, video) for e-learning, museums, and enterprise use cases; deployed across 5+ museums

- Built an internal framework to extract information from tables, spreadsheets, and enterprise documents using NLP and RAG; reduced table extraction errors by 75% vs open-source alternatives (pdfplumber, PyMuPDF), outperforming commercial tools including Azure Document Intelligence on structured specification tables
- Led a team of 2 junior engineers from pilot to production, owning the AI architecture and shipping 4 projects successfully

Machine Learning Engineer

Koexai

Feb 2023 – Nov 2023

Catania, Italy

- Built end-to-end ML pipelines on AWS (S3, Glue, Athena) for the largest coworking space in Catania (€1M revenue)
- Developed unsupervised ML models for customer segmentation (500+ customers) using K-Means, Hierarchical Clustering, and DBSCAN; enabled more granular price allocation across customer segments
- Automated recurring data entry processes, reducing operational time by 80% and freeing 2 employees for higher-value work

Writing, Contributions and Leadership

- Open-source contributor to **IBM/Agentic** (Python framework for agentic computation): identified and fixed a production bug in `PydanticTransducerVLLM`, resolved duplicate dependencies — PR merged into main (Jan 2026)
- Organized and technically led 2 **Hackatania GenAI hackathons**, defining challenge tracks, evaluating solutions, and mentoring participating teams
- Selected as participant in the GenAI Hackathon hosted by IBM & Datapizza in Milan (top 10% of 600+ applicants); built a Graph RAG project
- Authored technical articles on Retrieval-Augmented Generation and prompt engineering for Neodata Group

Education

M.Sc. in Data Science

University of Catania

2021 – 2023

Catania, Italy

Graduated 110/110 summa cum laude – GPA 4.0

Erasmus+ Mobility Semester in Artificial Intelligence

Goethe University Frankfurt

2023

Frankfurt, Germany

Bachelor in Business Management

University of Catania

2016 – 2020

Catania, Italy

Graduated 100/110 – GPA 3.3

Certifications

Introduction to LangGraph, LangChain (2024)

LangChain in Action: Develop LLM-Powered Applications (2024)

Intermediate SQL (2023)

Pre-Security Certificate, TryHackMe (2025)

Languages

Italian (native), English (professional), Spanish (professional)